

Homework 3

1. Operating System Concepts Chapter 1 Exercises: 1.14, 1.17, 1.19, 1.22 (20 points)

1.14 What is the purpose of interrupts? How does an interrupt differ from a trap? Can traps be generated intentionally by a user program? If so, for what purpose?

The purpose of interrupts throughout modern operating systems is to handle asynchronous events.

Differences:

1. The trap is a signal raised by a user program instructing the operating system to perform some functionality immediately. In contrast, the interrupt is a signal to the CPU emitted by hardware that indicates an event that requires immediate attention.
2. A trap also triggers OS functionality. It gives control to the trap handler. In contrast, an interrupt triggers the CPU to perform the interrupt handler routine.
3. A trap is synchronous and may occur after the execution of the instruction. In contrast, an interrupt is asynchronous and may occur at any time.
4. A trap is generated by a user program instruction. In contrast, the hardware devices generate an interrupt.

Yes, traps can be generated by a user program for instructing the operating system to do some instruction.

1.17 Some computer systems do not provide a privileged mode of operation in hardware. Is it possible to construct a secure operating system for these computer systems? Give arguments both that it is and that it is not possible.

If it is possible:

We can fix some system functions on the hardware and can't be rewritten. The security of the system depends on security of each application. If the programming languages are safe in memory space processing, it's possible to construct a secure operating system.

If it is not possible:

We can't promise all the applications are safe. If there isn't a privileged mode, malicious programs will rewrite the operating system code and may damage some vital functions of the operating system.

1.19 Rank the following storage systems from slowest to fastest:

a. Hard-disk drives b. Registers c. Optical disk d. Main memory e. Nonvolatile memory f. Magnetic tapes g. Cache

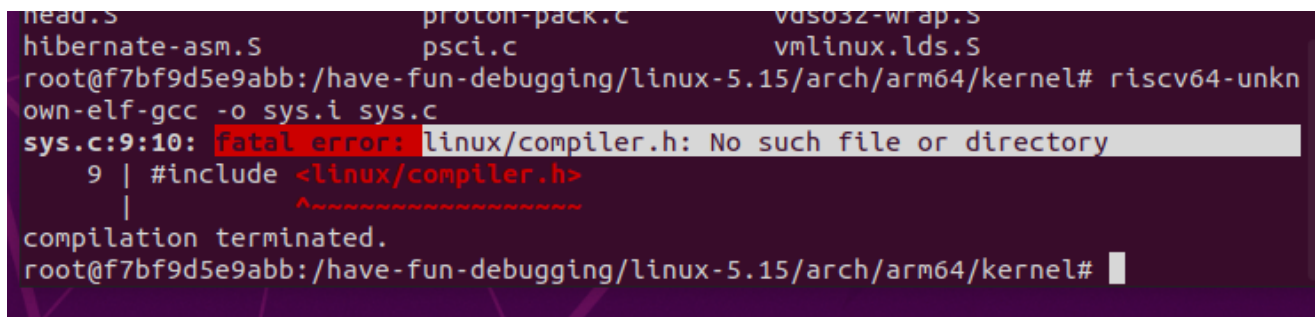
1. f. Magnetic tapes
2. c. Optical disk
3. a. Hard-disk drive
4. e. Nonvolatile memory
5. d. Main memory
6. g. Cache
7. b. Registers

1.22 Describe a mechanism for enforcing memory protection in order to prevent a program from modifying the memory associated with other programs

Virtual memory can be used as a memory protection tool, which stores the permissions and memory addresses of each process. Kernel and user modes are set in the operating system. The system checks the mode of each process before executing the instruction. When a process attempts to perform some operations beyond its mode, it will trigger the system exception function.

2. Detail your steps about how to get arch/arm64/kernel/sys.i (10 points)

```
1 //启动处于停止状态的容器
2 $docker start oslab1 //因为我的容器名称叫oslab1
3 $docker ps //查看正在运行的容器，这步用于确定容器正在运行，可以省略
4 $docker exec -it oslab1 bash //将终端连入docker容器(oslab1)
5 root@(随机码):/#
```



```
head.S          proton-pack.c          vds032-wrap.S
hibernate-asm.S  psci.c              vmlinux.lds.S
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/arch/arm64/kernel# riscv64-unknown-elf-gcc -o sys.i sys.c
sys.c:9:10: fatal error: linux/compiler.h: No such file or directory
   9 | #include <linux/compiler.h>
     |          ^
compilation terminated.
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/arch/arm64/kernel#
```

直接编译遇到错误，原因是arm架构下要用aarch tool chain。进行一下更新。

在linux-5.15文件夹路径下，依次输入以下指令

```
1 #apt-get update
2 #apt-get install gcc-aarch64-linux-gnu
3 #aarch64-linux-gnu-gcc -v //查看gcc版本，如果有版本信息，这说明添加成功
4 #make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- defconfig//生成配置
5 /*这里会输出
6 ***Default configuration is based on 'defconfig'
7 #
8 # configuration written to .config
9 #
10 以上这四行*/
11
12 #apt-get install libssl-dev
13 #make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- arch/arm64/kernel/sys.i -j $(nproc) //编译
14 //nproc内核参数,是系统上的最大进程数。使用多线程编译一般会耗费大量内存，如果 -j 选项导致内存耗尽
(out of memory)，请尝试调低线程数c'd，比如 -j4, -j8 等。
```

3. Find system call table of Linux v5.15 for ARM32, RISC-V(32 bit), RISC-V(64 bit), x86(32 bit), x86_64 (50 points)

List source code file, the whole system call table with macro expanded, screenshot every step.

```
1 进入oslab1容器的linux-5.15文件夹
2 #find / -name 'syscall*' //一定要加*不然会报错
```

```

rst
/have-fun-debugging/linux-5.15/arch/arm/include/asm/syscall.h
/have-fun-debugging/linux-5.15/arch/arm/tools/syscall.tbl
/have-fun-debugging/linux-5.15/arch/arm/tools/syscallnr.sh
/have-fun-debugging/linux-5.15/arch/x86/include/asm/syscall.h
/have-fun-debugging/linux-5.15/arch/x86/include/asm/syscalls.h
/have-fun-debugging/linux-5.15/arch/x86/include/asm/syscall_wrapper.h
/have-fun-debugging/linux-5.15/arch/x86/um/syscalls_64.c
/have-fun-debugging/linux-5.15/arch/x86/um/asm/syscall.h
/have-fun-debugging/linux-5.15/arch/x86/um/shared/sysdep/syscalls.h
/have-fun-debugging/linux-5.15/arch/x86/um/shared/sysdep/syscalls_64.h
/have-fun-debugging/linux-5.15/arch/x86/um/shared/sysdep/syscalls_32.h
/have-fun-debugging/linux-5.15/arch/x86/um/syscalls_32.c
/have-fun-debugging/linux-5.15/arch/x86/entry/syscalls
/have-fun-debugging/linux-5.15/arch/x86/entry/syscalls/syscall_32.tbl
/have-fun-debugging/linux-5.15/arch/x86/entry/syscalls/syscall_64.tbl
/have-fun-debugging/linux-5.15/arch/x86/entry/syscall_64.c
/have-fun-debugging/linux-5.15/arch/x86/entry/syscall_x32.c
/have-fun-debugging/linux-5.15/arch/x86/entry/syscall_32.c
/have-fun-debugging/linux-5.15/arch/arc/include/asm/syscall.h
/have-fun-debugging/linux-5.15/arch/arc/include/asm/syscalls.h
/have-fun-debugging/linux-5.15/arch/openrisc/include/asm/syscall.h
/have-fun-debugging/linux-5.15/arch/openrisc/include/asm/syscalls.h
/have-fun-debugging/linux-5.15/arch/mips/include/asm/syscall.h
/have-fun-debugging/linux-5.15/arch/mips/kernel/syscalls
/have-fun-debugging/linux-5.15/arch/mips/kernel/syscalls/syscall_n64.tbl
/have-fun-debugging/linux-5.15/arch/mips/kernel/syscalls/syscall_n32.tbl
/have-fun-debugging/linux-5.15/arch/mips/kernel/syscalls/syscallnr.sh
/have-fun-debugging/linux-5.15/arch/mips/kernel/syscalls/syscall_o32.tbl

```

然后我们得到了所有文件名包含syscall的文件，在一共166个文件中我们可以发现几个有用的。

```

1 - /have-fun-debugging/linux-5.15/arch/riscv/kernel/syscall_table.c
2 - /have-fun-debugging/linux-5.15/arch/x86/entry/syscalls/syscall_64.tbl
3 - /have-fun-debugging/linux-5.15/arch/x86/entry/syscalls/syscall_32.tbl
4 - /have-fun-debugging/linux-5.15/arch/arm/tools/syscall.tbl

```

.tbl文件可直接在terminal里通过cat指令查看，下面三个都可以这样看到system call table，.c文件则需要通过编译获得。以下为实验截图：

- /have-fun-debugging/linux-5.15/arch/arm/tools/syscall.tbl

```

root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15# cat arch/arm/tools/syscall.tbl
#
# Linux system call numbers and entry vectors
#
# The format is:
# <num> <abi> <name> [ <entry point> [ <oabi compat entry point>]]
#
# Where abi is:
# common - for system calls shared between oabi and eabi (may have compat)
# oabi - for oabi-only system calls (may have compat)
# eabi - for eabi-only system calls
#
# For each syscall number, "common" is mutually exclusive with oabi and eabi
#
0 common restart_syscall sys_restart_syscall
1 common exit sys_exit
2 common fork sys_fork
3 common read sys_read
4 common write sys_write
5 common open sys_open
6 common close sys_close
# 7 was sys_waitpid
8 common creat sys_creat
9 common link sys_link
10 common unlink sys_unlink
11 common execve sys_execve
12 common chdir sys_chdir
13 oabi time sys_time32
14 common mknod sys_mknod
15 common chmod sys_chmod
16 common lchown sys_lchown16
# 17 was sys_break
# 18 was sys_stat
19 common lseek sys_lseek
20 common getpid sys_getpid
21 common mount sys_mount
22 oabi umount sys_oldumount
23 common setuid sys_setuid16
24 common getuid sys_getuid16

```

```

# 18 was sys_stat
19 common lseek sys_lseek
20 common getpid sys_getpid
21 common mount sys_mount
22 oabi umount sys_oldumount
23 common setuid sys_setuid16
24 common getuid sys_getuid16
25 oabi stime sys_stime32
26 common ptrace sys_ptrace
27 oabi alarm sys_alarm
# 28 was sys_fstat
29 common pause sys_pause
30 oabi utime sys_utime32
# 31 was sys_stty
# 32 was sys_gtty
33 common access sys_access
34 common nice sys_nice
# 35 was sys_ftime
36 common sync sys_sync
37 common kill sys_kill
38 common rename sys_rename
39 common mkdir sys_mkdir
40 common rmdir sys_rmdir
41 common dup sys_dup
42 common pipe sys_pipe
43 common times sys_times
# 44 was sys_prof
45 common brk sys_brk
46 common setgid sys_setgid16
47 common getgid sys_getgid16
# 48 was sys_signal
49 common geteuid sys_geteuid16
50 common getegid sys_getegid16
51 common acct sys_acct
52 common umount2 sys_unmount
# 53 was sys_lock
54 common ioctl sys_ioctl
55 common fcntl sys_fcntl
# 56 was sys_mmap

```

```
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15

411 common timerfd_settime64 sys_timerfd_settime
412 common utimensat_time64 sys_utimensat
413 common pselect6_time64 sys_pselect6
414 common ppoll_time64 sys_ppoll
416 common io_pgetevents_time64 sys_io_pgetevents
417 common recvmmsg_time64 sys_recvmmsg
418 common mq_timedsend_time64 sys_mq_timedsend
419 common mq_timedreceive_time64 sys_mq_timedreceive
420 common sentimedop_time64 sys_sentimedop
421 common rt_sigtimedwait_time64 sys_rt_sigtimedwait
422 common futex_time64 sys_futex
423 common sched_rr_get_interval_time64 sys_sched_rr_get_interval
424 common pidfd_send_signal sys_pidfd_send_signal
425 common io_uring_setup sys_io_uring_setup
426 common io_uring_enter sys_io_uring_enter
427 common io_uring_register sys_io_uring_register
428 common open_tree sys_open_tree
429 common move_mount sys_move_mount
430 common fsopen sys_fsopen
431 common fsconfig sys_fsconfig
432 common fsmount sys_fsmount
433 common fspick sys_fspick
434 common pidfd_open sys_pidfd_open
435 common clone3 sys_clone3
436 common close_range sys_close_range
437 common openat2 sys_openat2
438 common pidfd_getfd sys_pidfd_getfd
439 common faccessat2 sys_faccessat2
440 common process_madvise sys_process_madvise
441 common epoll_pwait2 sys_epoll_pwait2
442 common mount_setattr sys_mount_setattr
443 common quotactl_fd sys_quotactl_fd
444 common landlock_create_ruleset sys_landlock_create_ruleset
445 common landlock_add_rule sys_landlock_add_rule
446 common landlock_restrict_self sys_landlock_restrict_self
# 447 reserved for memfd_secret
448 common process_mrelease sys_process_mrelease
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15#
```

■ /have-fun-debugging/linux-5.15/arch/x86/entry/syscalls/syscall_32.tbl

```
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15# cat arch/x86/entry/syscalls/syscall_32.tbl
#
# 32-bit system call numbers and entry vectors
#
# The format is:
# <number> <abi> <name> <entry point> <compat entry point>
#
# The __ia32_sys and __ia32_compat_sys stubs are created on-the-fly for
# sys_*() system calls and compat_sys_*() compat system calls if
# IA32_EMULATION is defined, and expect struct pt_regs *regs as their only
# parameter.
#
# The abi is always "i386" for this file.
#
0 i386 restart_syscall sys_restart_syscall
1 i386 exit sys_exit
2 i386 fork sys_fork
3 i386 read sys_read
4 i386 write sys_write
5 i386 open sys_open compat_sys_open
6 i386 close sys_close
7 i386 waitpid sys_waitpid
8 i386 creat sys_creat
9 i386 link sys_link
10 i386 unlink sys_unlink
11 i386 execve sys_execve compat_sys_execve
12 i386 chdir sys_chdir
13 i386 time sys_time32
14 i386 mknod sys_mknod
15 i386 chmod sys_chmod
16 i386 lchown sys_lchown16
17 i386 break
18 i386 oldstat sys_stat
19 i386 lseek sys_lseek compat_sys_lseek
20 i386 getpid sys_getpid
21 i386 mount sys_mount
22 i386 umount sys_oldumount
23 i386 setuid sys_setuid16
24 i386 getuid sys_getuid16
```

root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15				
70	i386	setreuid	sys_setreuid16	
71	i386	setregid	sys_setregid16	
72	i386	sigsuspend	sys_sigsuspend	
73	i386	sigpending	sys_sigpending	compat_sys_sigpending
74	i386	sethostname	sys_sethostname	
75	i386	setrlimit	sys_setrlimit	compat_sys_setrlimit
76	i386	getrlimit	sys_old_getrlimit	compat_sys_old_getrlimit
77	i386	getrusage	sys_getrusage	compat_sys_getrusage
78	i386	gettimeofday	sys_gettimeofday	compat_sys_gettimeofday
79	i386	settimeofday	sys_settimeofday	compat_sys_settimeofday
80	i386	getgroups	sys_getgroups16	
81	i386	setgroups	sys_setgroups16	
82	i386	select	sys_old_select	compat_sys_old_select
83	i386	symlink	sys_symlink	
84	i386	oldlstat	sys_lstat	
85	i386	readlink	sys_readlink	
86	i386	uselib	sys_uselib	
87	i386	swapon	sys_swapon	
88	i386	reboot	sys_reboot	
89	i386	readdir	sys_old_readdir	compat_sys_old_readdir
90	i386	mmap	sys_old_mmap	compat_sys_ia32_mmap
91	i386	munmap	sys_munmap	
92	i386	truncate	sys_truncate	compat_sys_truncate
93	i386	ftruncate	sys_ftruncate	compat_sys_ftruncate
94	i386	fchmod	sys_fchmod	
95	i386	fchown	sys_fchown16	
96	i386	getpriority	sys_getpriority	
97	i386	setpriority	sys_setpriority	
98	i386	profil		
99	i386	statfs	sys_statfs	compat_sys_statfs
100	i386	fstatfs	sys_fstatfs	compat_sys_fstatfs
101	i386	ioperm	sys_ioperm	
102	i386	socketcall	sys_socketcall	compat_sys_socketcall
103	i386	syslog	sys_syslog	
104	i386	setitimer	sys_setitimer	compat_sys_setitimer
105	i386	getitimer	sys_getitimer	compat_sys_getitimer
106	i386	stat	sys_newstat	compat_sys_newstat
107	i386	lstat	sys_newlstat	compat_sys_newlstat
108	i386	fstat	sys_newfstat	compat_sys_newfstat

root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15				
411	i386	timerfd_settime64	sys_timerfd_settime	
412	i386	utimensat_time64	sys_utimensat	
413	i386	pselect6_time64	sys_pselect6	compat_sys_pselect6_time64
414	i386	ppoll_time64	sys_ppoll	compat_sys_ppoll_time64
416	i386	io_pgetevents_time64	sys_io_pgetevents	
417	i386	recvmmsg_time64	sys_recvmmsg	compat_sys_recvmmsg_time64
418	i386	mq_timedsend_time64	sys_mq_timedsend	
419	i386	mq_timedreceive_time64	sys_mq_timedreceive	
420	i386	semtimedop_time64	sys_semtimedop	
421	i386	rt_sigtimedwait_time64	sys_rt_sigtimedwait	compat_sys_rt_sigtimedwait_time64
422	i386	futex_time64	sys_futex	
423	i386	sched_rr_get_interval_time64	sys_sched_rr_get_interval	
424	i386	pidfd_send_signal	sys_pidfd_send_signal	
425	i386	io_uring_setup	sys_io_uring_setup	
426	i386	io_uring_enter	sys_io_uring_enter	
427	i386	io_uring_register	sys_io_uring_register	
428	i386	open_tree	sys_open_tree	
429	i386	move_mount	sys_move_mount	
430	i386	fsopen	sys_fsopen	
431	i386	fsconfig	sys_fsconfig	
432	i386	fsmount	sys_fsmount	
433	i386	fspick	sys_fspick	
434	i386	pidfd_open	sys_pidfd_open	
435	i386	clone3	sys_clone3	
436	i386	close_range	sys_close_range	
437	i386	openat2	sys_openat2	
438	i386	pidfd_getfd	sys_pidfd_getfd	
439	i386	faccessat2	sys_faccessat2	
440	i386	process_madvise	sys_process_madvise	
441	i386	epoll_pwait2	sys_epoll_pwait2	compat_sys_epoll_pwait2
442	i386	mount_setattr	sys_mount_setattr	
443	i386	quotactl_fd	sys_quotactl_fd	
444	i386	landlock_create_ruleset	sys_landlock_create_ruleset	
445	i386	landlock_add_rule	sys_landlock_add_rule	
446	i386	landlock_restrict_self	sys_landlock_restrict_self	
447	i386	memfd_secret	sys_memfd_secret	
448	i386	process_mrelease	sys_process_mrelease	

共计448个系统调用函数表

- /have-fun-debugging/linux-5.15/arch/x86/entry/syscalls/syscall_64.tbl



```
448      i386      process_mrelease      sys_process_mrelease
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15# cat arch/x86/entry/syscalls/syscall_64.tbl
#
# 64-bit system call numbers and entry vectors
#
# The format is:
# <number> <abi> <name> <entry point>
#
# The __x64_sys_*() stubs are created on-the-fly for sys_*() system calls
#
# The abi is "common", "64" or "x32" for this file.
#
0      common    read                      sys_read
1      common    write                     sys_write
2      common    open                      sys_open
3      common    close                     sys_close
4      common    stat                      sys_newstat
5      common    fstat                     sys_newfstat
6      common    lstat                     sys_newlstat
7      common    poll                      sys_poll
8      common    lseek                     sys_lseek
9      common    mmap                      sys_mmap
10     common    mprotect                   sys_mprotect
11     common    munmap                    sys_munmap
12     common    brk                       sys_brk
13     64        rt_sigaction              sys_rt_sigaction
14     common    rt_sigprocmask            sys_rt_sigprocmask
15     64        rt_sigreturn              sys_rt_sigreturn
16     64        ioctl                     sys_ioctl
17     common    pread64                    sys_pread64
18     common    pwrite64                   sys_pwrite64
19     64        readv                      sys_readv
20     64        writev                     sys_writev
21     common    access                     sys_access
22     common    pipe                       sys_pipe
23     common    select                     sys_select
24     common    sched_yield                sys_sched_yield
25     common    mremap                     sys_mremap
```



```
332     common    statx                     sys_statx
333     common    io_pgetevents              sys_io_pgetevents
334     common    rseq                       sys_rseq
# don't use numbers 387 through 423, add new calls after the last
# 'common' entry
424     common    pidfd_send_signal          sys_pidfd_send_signal
425     common    io_uring_setup              sys_io_uring_setup
426     common    io_uring_enter              sys_io_uring_enter
427     common    io_uring_register            sys_io_uring_register
428     common    open_tree                    sys_open_tree
429     common    move_mount                   sys_move_mount
430     common    fsopen                       sys_fsopen
431     common    fsconfig                     sys_fsconfig
432     common    fsmount                       sys_fsmount
433     common    fspick                       sys_fspick
434     common    pidfd_open                   sys_pidfd_open
435     common    clone3                       sys_clone3
436     common    close_range                  sys_close_range
437     common    openat2                       sys_openat2
438     common    pidfd_getfd                  sys_pidfd_getfd
439     common    faccessat2                   sys_faccessat2
440     common    process_madvise              sys_process_madvise
441     common    epoll_pwait2                 sys_epoll_pwait2
442     common    mount_setattr                sys_mount_setattr
443     common    quotactl_fd                  sys_quotactl_fd
444     common    landlock_create_ruleset      sys_landlock_create_ruleset
445     common    landlock_add_rule            sys_landlock_add_rule
446     common    landlock_restrict_self       sys_landlock_restrict_self
447     common    memfd_secret                 sys_memfd_secret
448     common    process_mrelease             sys_process_mrelease
#
# Due to a historical design error, certain syscalls are numbered differently
# in x32 as compared to native x86_64. These syscalls have numbers 512-547.
# Do not add new syscalls to this range. Numbers 548 and above are available
# for non-x32 use.
#
512     x32       rt_sigaction                compat_sys_rt_sigaction
```

```
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15

513 x32 rt_sigreturn compat_sys_x32_rt_sigreturn
514 x32 ioctl compat_sys_ioctl
515 x32 readv sys_readv
516 x32 writev sys_writev
517 x32 recvfrom compat_sys_recvfrom
518 x32 sendmsg compat_sys_sendmsg
519 x32 recvmmsg compat_sys_recvmmsg
520 x32 execve compat_sys_execve
521 x32 ptrace compat_sys_ptrace
522 x32 rt_sigpending compat_sys_rt_sigpending
523 x32 rt_sigtimedwait compat_sys_rt_sigtimedwait_time64
524 x32 rt_sigqueueinfo compat_sys_rt_sigqueueinfo
525 x32 sigaltstack compat_sys_sigaltstack
526 x32 timer_create compat_sys_timer_create
527 x32 mq_notify compat_sys_mq_notify
528 x32 kexec_load compat_sys_kexec_load
529 x32 waitid compat_sys_waitid
530 x32 set_robust_list compat_sys_set_robust_list
531 x32 get_robust_list compat_sys_get_robust_list
532 x32 vmsplce sys_vmsplce
533 x32 move_pages sys_move_pages
534 x32 preadv compat_sys_preadv64
535 x32 pwritev compat_sys_pwritev64
536 x32 rt_tgsigqueueinfo compat_sys_rt_tgsigqueueinfo
537 x32 recvmmsg compat_sys_recvmmsg_time64
538 x32 sendmmsg compat_sys_sendmmsg
539 x32 process_vm_readv sys_process_vm_readv
540 x32 process_vm_writev sys_process_vm_writev
541 x32 setsockopt sys_setsockopt
542 x32 getsockopt sys_getsockopt
543 x32 io_setup compat_sys_io_setup
544 x32 io_submit compat_sys_io_submit
545 x32 execveat compat_sys_execveat
546 x32 preadv2 compat_sys_preadv64v2
547 x32 pwritev2 compat_sys_pwritev64v2
# This is the end of the legacy x32 range. Numbers 548 and above are
# not special and are not to be used for x32-specific syscalls.
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15#
```

共计547个系统调用函数

- /have-fun-debugging/linux-5.15/arch/riscv/kernel/syscall_table.c

在riscv文件夹中，只有这一个syscall_table.c函数记录了系统函数调用表。在64位和32位不同编译器中编出的汇编语言可能不一样，但其函数调用表的序号和函数是一样的，故在此只展示出64位riscv64-unknown-linux-gnu-gcc编译出的系统调用表。

- ```
1 #make ARCH=riscv CROSS_COMPILE=riscv64-unknown-linux-gnu- arch/riscv/kernel/syscall_table.i -
 j$(nproc)
2 这里要注意中间是.i不是.c，中间是目标文件
```

```
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15

void * const sys_call_table[449] = {
 [0 ... 449 - 1] = sys_ni_syscall,
 # 1 "/arch/riscv/include/asm/unistd.h" 1
 # 14 "/arch/riscv/include/asm/unistd.h"
 # 1 "/arch/riscv/include/uapi/asm/unistd.h" 1
 # 25 "/arch/riscv/include/uapi/asm/unistd.h"
 # 1 "/include/uapi/asm-generic/unistd.h" 1
 # 34 "/include/uapi/asm-generic/unistd.h"
 [0] = (sys_io_setup),

 [1] = (sys_io_destroy),

 [2] = (sys_io_submit),

 [3] = (sys_io_cancel),

 [4] = (sys_io_getevents),

 [5] = (sys_setxattr),

 [6] = (sys_lsetxattr),

 [7] = (sys_fsetxattr),

 [8] = (sys_getxattr),
```



```
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15

[440] = (sys_process_madvise),
[441] = (sys_epoll_pwait2),
[442] = (sys_mount_setattr),
[443] = (sys_quotactl_fd),

[444] = (sys_landlock_create_ruleset),
[445] = (sys_landlock_add_rule),
[446] = (sys_landlock_restrict_self),

[447] = (sys_memfd_secret),

[448] = (sys_process_mrelease),
26 "/arch/riscv/include/uapi/asm/unistd.h" 2
44 "/arch/riscv/include/uapi/asm/unistd.h"
[(244 + 15)] = (sys_riscv_flush_icache),
15 "/arch/riscv/include/asm/unistd.h" 2
18 "arch/riscv/kernel/syscall_table.c" 2
};
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15# S
```

#### 4.Explain what is ELF file? Try readelf and objdump command on an ELF file, give screenshot of the output.

Run an ELF file and cat /proc/PID/maps to give its memory layout.

Executable and Linkable Format. It is a common standard file format for executable files, object code, shared libraries, and core dumps.

```
1 #touch test.c //新建test.c
2 #vi test.c //进入编辑test.c
3 Esc + :exit //退出编辑
4 #gcc test.c //没有规定文件名称，默认生成了a.out
5 #readelf -a a.out //all显示全部信息，等价于-h-l-S-s-r-d-V-A-I
```

#### readelf 只能看elf文件的信息

- 选项-h(elfheader)，显示elf文件开始的文件头信息。后面文章会补上具体说明。
- 选项-l(programheaders)，segments显示程序头（段头）信息(如果有数据的话)。后面文章会补上具体说明。
- 选项-S(sectionheaders)，sections显示节头信息(如果有数据的话)。后面文章会补上具体说明。
- 选项-a，all显示全部信息，等价于-h-l-S-s-r-d-V-A-I。
- 选项-g(sectiongroups)，显示节组信息(如果有数据的话)。
- 选项-t，section-details显示节的详细信息(-S的)。
- 选项-s，symbols显示符号表段中的项（如果有数据的话）。
- 选项-e，headers显示全部头信息，等价于:-h-l-S。
- 选项-n，notes显示note段（内核注释）的信息。
- 选项-r，relocs显示可重定位段的信息。
- 选项-u，unwind显示unwind段信息。当前只支持IA64ELF的unwind段信息。
- 选项-d，dynamic显示动态段的信息。
- 选项-V，version-info显示版本段的信息。
- 选项-A，arch-specific显示CPU构架信息。
- 选项-I，histogram显示符号的时候，显示bucketlist长度的柱状图。
- 选项-x,hex-dump=以16进制方式显示指定段内内容。number指定段表中段的索引，或字符串指定文件中的段名
- 选项-D，use-dynamic使用动态段中的符号表显示符号，而不是使用符号段。

- 选项-v, version显示readelf的版本信息。
- 选项-H, help显示readelf所支持的命令行选项。

```

root@f7bf9d5e9abb: /h
a.out test.c
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/newtest# readelf -a a.out
ELF Header:
 Magic: 7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
 Class: ELF64
 Data: 2's complement, little endian
 Version: 1 (current)
 OS/ABI: UNIX - System V
 ABI Version: 0
 Type: DYN (Shared object file)
 Machine: Advanced Micro Devices X86-64
 Version: 0x1
 Entry point address: 0x1060
 Start of program headers: 64 (bytes into file)
 Start of section headers: 14712 (bytes into file)
 Flags: 0x0
 Size of this header: 64 (bytes)
 Size of program headers: 56 (bytes)
 Number of program headers: 13
 Size of section headers: 64 (bytes)
 Number of section headers: 31
 Section header string table index: 30

Section Headers:
[Nr] Name Type Address Offset
 Size EntSize Flags Link Info Align
[0] NULL 0000000000000000 00000000
 0000000000000000 0000000000000000 0 0 0
[1] .interp PROGBITS 0000000000000318 00000318
 000000000000001c 0000000000000000 A 0 0 1
[2] .note.gnu.property NOTE 0000000000000338 00000338
 0000000000000020 0000000000000000 A 0 0 8
[3] .note.gnu.build-id NOTE 0000000000000358 00000358
 0000000000000024 0000000000000000 A 0 0 4
[4] .note.ABI-tag NOTE 000000000000037c 0000037c
 0000000000000020 0000000000000000 A 0 0 4
[5] .gnu.hash GNU_HASH 00000000000003a0 000003a0
 0000000000000024 0000000000000000 A 6 0 8

```

```
root@f7bf9d5e9abb: /have-fu

[29] .strtab STRTAB 0000000000000000 00003658
0000000000000204 0000000000000000 0 0 1
[30] .shstrtab STRTAB 0000000000000000 0000385c
000000000000011a 0000000000000000 0 0 1
Key to Flags:
W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
l (large), p (processor specific)

There are no section groups in this file.

Program Headers:
Type Offset VirtAddr PhysAddr
 FileSiz MemSiz Flags Align
PHDR 0x0000000000000040 0x0000000000000040 0x0000000000000040
 0x00000000000002d8 0x00000000000002d8 R 0x8
INTERP 0x0000000000000318 0x0000000000000318 0x0000000000000318
 0x000000000000001c 0x000000000000001c R 0x1
 [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
LOAD 0x0000000000000000 0x0000000000000000 0x0000000000000000
 0x0000000000000600 0x0000000000000600 R 0x1000
LOAD 0x0000000000000100 0x0000000000000100 0x0000000000000100
 0x00000000000001f5 0x00000000000001f5 R E 0x1000
LOAD 0x0000000000000200 0x0000000000000200 0x0000000000000200
 0x0000000000000160 0x0000000000000160 R 0x1000
LOAD 0x00000000000002db8 0x00000000000003db8 0x00000000000003db8
 0x0000000000000258 0x0000000000000260 RW 0x1000
DYNAMIC 0x00000000000002dc8 0x00000000000003dc8 0x00000000000003dc8
 0x00000000000001f0 0x00000000000001f0 RW 0x8
NOTE 0x0000000000000338 0x0000000000000338 0x0000000000000338
 0x0000000000000020 0x0000000000000020 R 0x8
NOTE 0x0000000000000358 0x0000000000000358 0x0000000000000358
 0x0000000000000044 0x0000000000000044 R 0x4
GNU_PROPERTY 0x0000000000000338 0x0000000000000338 0x0000000000000338
 0x0000000000000020 0x0000000000000020 R 0x8
GNU_EH_FRAME 0x00000000000002014 0x00000000000002014 0x00000000000002014
 0x0000000000000044 0x0000000000000044 R 0x4
GNU_STACK 0x0000000000000000 0x0000000000000000 0x0000000000000000
```

```
root@f7bf9d5e9abb: /have-fu

11
12 .init_array .fini_array .dynamic .got

Dynamic section at offset 0x2dc8 contains 27 entries:
Tag Type Name/Value
0x0000000000000001 (NEEDED) Shared library: [libc.so.6]
0x000000000000000c (INIT) 0x1000
0x000000000000000d (FINI) 0x11e8
0x0000000000000019 (INIT_ARRAY) 0x3db8
0x000000000000001b (INIT_ARRAYSZ) 8 (bytes)
0x000000000000001a (FINI_ARRAY) 0x3dc0
0x000000000000001c (FINI_ARRAYSZ) 8 (bytes)
0x000000006ffffef5 (GNU_HASH) 0x3a0
0x0000000000000005 (STRTAB) 0x470
0x0000000000000006 (SYMTAB) 0x3c8
0x000000000000000a (STRSZ) 132 (bytes)
0x000000000000000b (SYMENT) 24 (bytes)
0x0000000000000015 (DEBUG) 0x0
0x0000000000000003 (PLTGOT) 0x3fb8
0x0000000000000002 (PLTRELSZ) 24 (bytes)
0x0000000000000014 (PLTREL) RELA
0x0000000000000017 (JMPREL) 0x5e8
0x0000000000000007 (RELA) 0x528
0x0000000000000008 (RELASZ) 192 (bytes)
0x0000000000000009 (RELAENT) 24 (bytes)
0x000000000000001e (FLAGS) BIND_NOW
0x000000006ffffffb (FLAGS_1) Flags: NOW PIE
0x000000006ffffffe (VERNEED) 0x508
0x000000006fffffff (VERNEEDNUM) 1
0x000000006ffffff0 (VERSYM) 0x4f4
0x000000006ffffff9 (RELACOUNT) 3
0x0000000000000000 (NULL) 0x0

Relocation section '.rela.dyn' at offset 0x528 contains 8 entries:
Offset Info Type Sym. Value Sym. Name + Addend
000000003db8 000000000008 R_X86_64_RELATIVE 1140
000000003dc0 000000000008 R_X86_64_RELATIVE 1100
000000004008 000000000008 R_X86_64_RELATIVE 4008
```

(后信息省略)

## objdump

```
1 objdump -d a.out // -d 参数看代码段反汇编结果
2 objdump -t a.out // 显示文件的符号表入口。
3 objdump -t libc.a grep -w printf // 查找 printf 在 libc.a 库的哪个目标文件
4 objdump -h simple.o // 显示目标文件各个section的头部摘要信息
5 objdump -r simple.o // 查看重定位表
6 objdump -f simple.o // 显示objfile中每个文件的整体头部摘要信息
7 objdump -s simple.o // 显示指定section的完整内容
8 objdump -x simple.o // 显示所可用的头信息
9 objdump -a simple.o // 显示档案库的成员信息
```

反汇编结果如下

```
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15/newtest
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/newtest# objdump -d a.o
objdump: 'a.o': No such file
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/newtest# objdump -d a.out
a.out: file format elf64-x86-64

Disassembly of section .init:

0000000000001000 <.init>:
 1000: f3 0f 1e fa endbr64
 1004: 48 83 ec 08 sub $0x8,%rsp
 1008: 48 8b 05 d9 2f 00 00 mov 0x2fd9(%rip),%rax # 3fe8 <__gmon_start__>
 100f: 48 85 c0 test %rax,%rax
 1012: 74 02 je 1016 <_init+0x16>
 1014: ff d0 callq *%rax
 1016: 48 83 c4 08 add $0x8,%rsp
 101a: c3 retq

Disassembly of section .plt:

0000000000001020 <.plt>:
 1020: ff 35 9a 2f 00 00 pushq 0x2f9a(%rip) # 3fc0 <_GLOBAL_OFFSET_TABLE_+0x8>
 1026: f2 ff 25 9b 2f 00 00 bnd jmpq *0x2f9b(%rip) # 3fc8 <_GLOBAL_OFFSET_TABLE_+0x10>
 102d: 0f 1f 00 nopl (%rax)
 1030: f3 0f 1e fa endbr64
 1034: 68 00 00 00 00 00 pushq $0x0
 1039: f2 e9 e1 ff ff ff bnd jmpq 1020 <.plt>
 103f: 90 nop

Disassembly of section .plt.got:

0000000000001040 <__cxa_finalize@plt>:
 1040: f3 0f 1e fa endbr64
 1044: f2 ff 25 ad 2f 00 00 bnd jmpq *0x2fad(%rip) # 3ff8 <__cxa_finalize@GLIBC_2.2.5>
 104b: 0f 1f 44 00 00 00 nopl 0x0(%rax,%rax,1)

Disassembly of section .plt.sec:
```

符号表入口

```
root@f7bf9d5e9abb: /have-fun-debugging/linux-5.15/newtest

11f4: c3 retq
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/newtest#
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/newtest# objdump -t a.out

a.out: file format elf64-x86-64

SYMBOL TABLE:
0000000000000318 l d .interp 0000000000000000 .interp
0000000000000338 l d .note.gnu.property 0000000000000000 .note.gnu.property
0000000000000358 l d .note.gnu.build-id 0000000000000000 .note.gnu.build-id
000000000000037c l d .note.ABI-tag 0000000000000000 .note.ABI-tag
00000000000003a0 l d .gnu.hash 0000000000000000 .gnu.hash
00000000000003c8 l d .dynsym 0000000000000000 .dynsym
0000000000000470 l d .dynstr 0000000000000000 .dynstr
00000000000004f4 l d .gnu.version 0000000000000000 .gnu.version
0000000000000508 l d .gnu.version_r 0000000000000000 .gnu.version_r
0000000000000528 l d .rela.dyn 0000000000000000 .rela.dyn
00000000000005e8 l d .rela.plt 0000000000000000 .rela.plt
0000000000001000 l d .init 0000000000000000 .init
0000000000001020 l d .plt 0000000000000000 .plt
0000000000001040 l d .plt.got 0000000000000000 .plt.got
0000000000001050 l d .plt.sec 0000000000000000 .plt.sec
0000000000001060 l d .text 0000000000000000 .text
00000000000011e8 l d .fini 0000000000000000 .fini
0000000000002000 l d .rodata 0000000000000000 .rodata
0000000000002014 l d .eh_frame_hdr 0000000000000000 .eh_frame_hdr
0000000000002058 l d .eh_frame 0000000000000000 .eh_frame
0000000000003db8 l d .init_array 0000000000000000 .init_array
0000000000003dc0 l d .fini_array 0000000000000000 .fini_array
0000000000003dc8 l d .dynamic 0000000000000000 .dynamic
0000000000003fb8 l d .got 0000000000000000 .got
0000000000004000 l d .data 0000000000000000 .data
0000000000004010 l d .bss 0000000000000000 .bss
0000000000000000 l d .comment 0000000000000000 .comment
0000000000000000 l df *ABS* 0000000000000000 crtstuff.c
0000000000001090 l F .text 0000000000000000 deregister_tm_clones
00000000000010c0 l F .text 0000000000000000 register_tm_clones
0000000000001100 l F .text 0000000000000000 __do_global_ctors_aux
0000000000001010 l F .bss 0000000000000000 completed_ctors
```

## Run an ELF file and cat /proc/PID/maps to give its memory layout.

```
/have-fun-debugging/linux-5.15/newtest# ps aux | grep pmap
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15# ps aux | grep pmap
root 3082 0.0 0.0 3312 732 pts/2 S+ 09:29 0:00 grep --color=auto pmap
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15#
```

1 | #ps aux |grep pmap //查看程序pid

可以看到程序pmap的pid为3082，但是这个程序运行结束后，它所分配的内存也被回收了，所以在下一步搜索中并看不到/3082/maps。此时文件夹中并没有/proc路径，但是经过搜索我们可以看到这个路径是在系统根目录下的。

```
cat: /proc/3085/maps: No such file or directory
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/newtest# find / -name 'maps*'
/proc/1/task/1/maps
/proc/1/maps
/proc/34/task/34/maps
/proc/34/maps
/proc/3039/task/3039/maps
/proc/3039/maps
/proc/3087/task/3087/maps
/proc/3087/maps
/have-fun-debugging/linux-5.15/drivers/mtd/maps
/have-fun-debugging/linux-5.15/tools/perf/tests/maps.c
/have-fun-debugging/linux-5.15/tools/perf/util/maps.h
```

通过查找可以发现有这么几个进程。其中3087为terminal窗口执行每一条指令的进程，每输入一条指令该数都会增加，查询指令对应的进程号是3087，在此次查询结束后是无法通过cat /proc/3087/maps访问到的。故选择进程3039进行查看。

可以看到其中每个进程运行的：

- address: 0085d000-00872000 虚拟内存区域的起始和终止地址文件所占的地址空间
- perms:rw-p 权限：r=read, w=write, x=execute, s=shared, p=private(copy on write)
- offset: 00000000 虚拟内存区域在被映射文件中的偏移量
- dev: 03:08 文件的主设备号和次设备号



- inode: 设备的节点号, 0表示没有节点与内存相对应
- name: /lib/ld-2.2.12.so 被映射文件的文件名

```

/have-fun-debugging/linux-5.15/boots/perf/dec/maps.m
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/newtest# cat /proc/3039/m
aps
55e17a5eb000-55e17a618000 r--p 00000000 00:34 674295 /
usr/bin/bash
55e17a618000-55e17a6c9000 r-xp 0002d000 00:34 674295 /
usr/bin/bash
55e17a6c9000-55e17a700000 r--p 000de000 00:34 674295 /
usr/bin/bash
55e17a700000-55e17a704000 r--p 00114000 00:34 674295 /
usr/bin/bash
55e17a704000-55e17a70d000 rw-p 00118000 00:34 674295 /
usr/bin/bash
55e17a70d000-55e17a717000 rw-p 00000000 00:00 0
55e17ab62000-55e17abc5000 rw-p 00000000 00:00 0 [
heap]
7f55ae2fc000-7f55ae2ff000 r--p 00000000 00:34 679291 /
usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
7f55ae2ff000-7f55ae306000 r-xp 00003000 00:34 679291 /
usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
7f55ae306000-7f55ae308000 r--p 0000a000 00:34 679291 /
usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f55ae53a000-7f55ae53b000 rw-p 0002e000 00:34 679423 /
usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f55ae53b000-7f55ae53d000 rw-p 00000000 00:00 0
7f55ae542000-7f55ae543000 r--p 00000000 00:34 678884 /
usr/lib/x86_64-linux-gnu/ld-2.31.so
7f55ae543000-7f55ae566000 r-xp 00001000 00:34 678884 /
usr/lib/x86_64-linux-gnu/ld-2.31.so
7f55ae566000-7f55ae56e000 r--p 00024000 00:34 678884 /
usr/lib/x86_64-linux-gnu/ld-2.31.so
7f55ae56f000-7f55ae570000 r--p 0002c000 00:34 678884 /
usr/lib/x86_64-linux-gnu/ld-2.31.so
7f55ae570000-7f55ae571000 rw-p 0002d000 00:34 678884 /
usr/lib/x86_64-linux-gnu/ld-2.31.so
7f55ae571000-7f55ae572000 rw-p 00000000 00:00 0
7ffefa6ab000-7ffefa6cc000 rw-p 00000000 00:00 0 [
stack]
7ffefa75d000-7ffefa761000 r--p 00000000 00:00 0 [
vvar]
7ffefa761000-7ffefa763000 r-xp 00000000 00:00 0 [
vdso]
ffffffffffff600000-ffffffffffff601000 --xp 00000000 00:00 0 [
vsyscall]
root@f7bf9d5e9abb:/have-fun-debugging/linux-5.15/newtest#

```

2, 3, 4 need to have screenshots.